

INDO-SWISS HACKATHON 2026 | ROUND 2 SUBMISSION

Supply Chain Decision Intelligence

Calibrated Uncertainty Quantification for Supply Chain Risk Triage

Category:	Solution & Technical White Paper
Authors:	Indo-Swiss Hackathon Team
Date:	June 30, 2026
Version:	1.0.0
Live Dashboard:	cognifyaiprediction.streamlit.app

1. Problem Understanding

The Business Problem

Modern retail supply chains are fundamentally broken in how they handle uncertainty. Planners rely almost exclusively on "point forecasts" — single-number predictions (e.g., "we will sell 42 units next week"). However, these point forecasts are virtually guaranteed to be wrong. When demand volatility hits (due to seasonality, promotions, or supply shocks), the supply chain reacts inefficiently, resulting in either massive stockouts (lost revenue) or bloated safety stocks (wasted capital).

Why Traditional Forecasting is Insufficient

Traditional machine learning (like raw XGBoost or ARIMA) outputs a deterministic value minimizing a loss function (like RMSE or MAPE). It provides no statistical indication of its own confidence. If a model predicts 10 units for a stable product and 10 units for a highly volatile product, traditional systems treat these forecasts identically. This blind spot forces businesses to apply blanket, heuristic-based safety stock rules (e.g., "always keep 20% extra inventory") which are systematically inefficient.

The Power of Calibrated Uncertainty

Instead of a single number, our system outputs a **Prediction Interval** (e.g., "we are 90% confident demand will fall between 8 and 15 units"). But simply producing an interval isn't enough; it must be *calibrated*. Calibration means that a 90% confidence interval empirically contains the true demand exactly 90% of the time, regardless of whether the product is highly seasonal, newly introduced, or consistently stable. Calibrated uncertainty is the mathematical foundation required to optimally balance the financial trade-off between holding cost (overstocking) and penalty cost (stockouts).

2. Proposed Solution

We have built an end-to-end Decision Intelligence Platform that bridges the gap between raw machine learning output and actionable business strategy.

Our workflow ingests raw retail sales data, generates a highly optimized gradient-boosted point forecast, and wraps it in a statistically guaranteed uncertainty interval using Split Conformal Prediction. A specialized Risk Engine then scores each product based on its uncertainty profile, and an Inventory Simulator determines the exact financial cost/benefit of intervening. Finally, everything is surfaced to business users through an intuitive, interactive Streamlit Executive Dashboard.

3. Technical Approach

Stage 1: Dataset Understanding

We utilized the M5 Forecasting dataset (Walmart hierarchical sales data) encompassing 3,049 products across 10 stores over 5 years. Initial EDA identified zero-inflated demand distributions, strong weekly/yearly seasonality, and high volatility clustering around promotional events and SNAP benefit days.

Stage 2: Feature Engineering

We constructed a rich feature space to capture non-linear demand drivers:

- **Temporal:** Day of week, month, year, event proximity.
- **Lag/Rolling:** 7-day, 14-day, 28-day rolling means and standard deviations to capture momentum and volatility.
- **Categorical:** Target encoding for items, departments, and stores.
- **Sparsity Handling:** Engineered a binary `is_zero_sales` flag to assist the model with intermittent demand.

Stage 3: Forecasting (XGBoost)




Figure 1: XGBoost Tweedie Point Forecast vs. Actual Walmart Sales (M5 Dataset)

We trained an XGBoost regressor using the Tweedie objective function, which is explicitly designed for zero-inflated, right-skewed data. Hyperparameters were rigorously optimized using Optuna, resulting in a model that captured the complex hierarchical interactions across 30,000+ time series.

Stage 4: Uncertainty Quantification

Quantile Regression vs Split Conformal Prediction Method Comparison — 90% Prediction Intervals




Figure 2: Prediction Interval Comparison: Pinball Loss Quantiles vs. Split Conformal Bounds

To generate prediction intervals, we implemented two distinct methodologies:

- 1. Quantile Regression:** Training separate models for the 5th and 95th percentiles using Pinball Loss.
- 2. Split Conformal Prediction:** A distribution-free, model-agnostic approach that calibrates intervals based on the empirical distribution of holdout residuals.

Decision: We ultimately selected Conformal Prediction for the final engine due to its superior marginal coverage guarantees and robustness against quantile crossing.

Stage 5: Calibration Evaluation




Figure 3: Split Conformal Calibration Curve showing nominal vs. empirical coverage

We rigorously evaluated the intervals using reliability diagrams and coverage analysis across different volatility segments. The Conformal Prediction intervals achieved ~87.5% empirical coverage (closely tracking the 90% target), proving that the model's confidence estimates are statistically reliable and safe for downstream financial optimization.

Stage 6: Risk Engine




Figure 4: Portfolio Risk Segmentation (Low, Medium, High) computed by the Risk Engine

We developed a proprietary composite Risk Score (0-100) based on four factors:

1. Interval Width (Model uncertainty)
2. Historical Forecast Error (Past inaccuracy)
3. Rolling Volatility (Inherent demand instability)
4. Calibration Penalty (Local model distrust)

Products are then segmented into Low (Green), Medium (Yellow), and High (Red) risk tiers.

Stage 7: Inventory Simulation




Figure 5: Operational Cost Comparison: Point-Forecast Ordering vs. Conformal-Risk Safety Stock

We built a Business Impact Simulator to compare two operational strategies:

- **Baseline System:** Order exactly the point forecast.
- **Intelligent System:** Order at the 90th percentile (upper bound) of the conformal interval for high-risk items.

The simulation applied real-world holding cost rates (5%) and stockout penalties (40%), proving that the risk-aware strategy drastically reduces stockouts and increases overall profitability.

Stage 8: Executive Dashboard

The entire ML pipeline and decision engine are packaged into a streamlined, four-page Streamlit application. Rather than simply showing data plots, the dashboard translates complex statistical models into plain-English recommendations ("Increase safety stock by 15 units. Penalty Mitigation Ratio: 5.6x") and includes simulated interactive workflows to push these parameters directly to a mock ERP.

4. Prototype Design

Dashboard & User Workflow




Figure 6: Executive Dashboard visualization showing average risk scores across product stores

The platform is designed as an interactive, zero-latency decision tool composed of exactly four sections:

- **1. Executive Overview:** Provides the "C-Suite" view (net savings, stockout reduction, service levels, avg risk) alongside risk distribution charts, an executive briefing, and a technical appendix explaining key assumptions.
- **2. Risk & Operations:** The operational workspace. Planners can triage high-risk items, select SKUs to view actual vs. forecast charts (with 90% conformal prediction bands), read plain-English risk factors, and click to simulate ERP actions (Approve PO, Send to ERP).
- **3. Scenario Simulator:** An interactive supply chain physics sandbox. Planners can adjust holding costs, stockout penalties, target confidence, replenishment lead time, and macro demand volatility to immediately recalculate optimal buffer bounds live.
- **4. Technical Engine:** A model validator for data scientists. Exposes Tweedie XGBoost error metrics (WAPE, RMSE), residuals distributions, conformal calibration curves, segment coverage breakdowns, and global feature importance.

- **Guided Demo Mode:** A global sidebar controller that walks judges through the real pages of the dashboard with narrative banners in under 3 minutes.
-

5. Feasibility Analysis

- **Deployment Feasibility:** The architecture is stateless and containerizable (Docker/Kubernetes ready). The dashboard (Streamlit) runs independently of the heavy ML training pipeline.
 - **Scalability:** Parquet data formats and Optuna-optimized XGBoost ensure the system scales efficiently. Inferencing and Conformal calibration are exceptionally fast ($O(1)$ relative to training time).
 - **Maintenance:** Split Conformal Prediction requires zero retraining to update interval widths; it only requires a fast recalibration step on recent residuals, drastically reducing compute costs.
 - **Infrastructure:** Can be seamlessly deployed on AWS EC2/SageMaker or GCP Cloud Run.
-

6. Expected Impact

Based on our simulation across 5,600 product-store testing periods:

- **Total Operating Costs:** Reduced by **15.6%**.
 - **Service Levels:** Improved from **83.3%** to **92.2%** (an 8.9 percentage point increase).
 - **Stockouts:** Eliminated **1,545 lost-sale units**.
 - **Financial ROI:** For every \$1 invested in targeted safety stock holding costs, the business avoids \$0.56 in stockout penalties, resulting in an overwhelmingly positive net financial return.
 - **Decision Quality:** Supply chain planners no longer guess safety stocks; they act on statistically guaranteed risk bounds.
-

7. Future Scope

1. **Multi-Echelon Optimization:** Extending the uncertainty propagation upstream to warehouses and distribution centers.
2. **Reinforcement Learning (RL):** Training an RL agent to continuously adjust the confidence level parameter (alpha) based on dynamic supply chain macro-economics.
3. **LLM Copilot:** Integrating an LLM (e.g., Gemini 1.5 Pro) directly into the dashboard to allow planners to query their supply chain risks using natural language ("Show me all high-risk dairy items in California").
4. **ERP Integration:** Deploying the inference API directly into SAP or Oracle to automate PO (Purchase Order) generation based on the Risk Engine's outputs.